

```

import os
import numpy as np
import tensorflow as tf
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D,
UpSampling2D, Concatenate, Dropout
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.metrics import MeanIoU, Precision, Recall
from tensorflow.keras.callbacks import ModelCheckpoint,
EarlyStopping
from tensorflow.keras.regularizers import l2
from tensorflow.keras.losses import CategoricalCrossentropy
import matplotlib.pyplot as plt

# Directorios con imágenes y máscaras ya separadas
train_image_dir = r"C:\Users\ismae\Desktop\TFG\Python
Windows\frames"
train_mask_dir =
r"C:\Users\ismae\Desktop\TFG\Mascaras\entrenamiento"
val_image_dir = r"C:\Users\ismae\Desktop\TFG\Python
Windows\frames"
val_mask_dir = r"C:\Users\ismae\Desktop\TFG\Mascaras\validacion"

# Parámetros del modelo
IMAGE_SIZE = (256, 256) # Tamaño al que redimensionaremos las
imágenes
NUM_CLASSES = 4 # Número de clases (Fondo, hueso, tendón,
límites)
BATCH_SIZE = 4
EPOCHS = 50

def desescalar_masks(masks):
    return (masks / 85).astype(np.uint8)

# Función para cargar imágenes y máscaras
def load_images_and_masks(image_dir, mask_dir, image_size):
    images = []
    masks = []
    for file_name in os.listdir(image_dir):
        image_path = os.path.join(image_dir, file_name)
        mask_path = os.path.join(mask_dir,
f"{os.path.splitext(file_name)[0]}_mask.png")

        if os.path.exists(image_path) and
os.path.exists(mask_path):
            # Cargar y redimensionar la imagen
            image =
tf.image.decode_image(tf.io.read_file(image_path), channels=3)
            image = tf.image.resize(image, image_size)

            # Cargar y redimensionar la máscara
            mask =
tf.image.decode_image(tf.io.read_file(mask_path), channels=1)
            mask = tf.image.resize(mask, image_size,
method='nearest') # Mantener valores discretos

```

```

        # Desescalar las máscaras (85, 170, 255 → 1, 2, 3)
        mask = desescalar_masks(mask.numpy())

        # Agregar a las listas
        images.append(image.numpy())
        masks.append(mask)

    images = np.array(images) / 255.0 # Normalizar imágenes (0-
1)
    masks = np.array(masks) # Asegurar que las máscaras sean
enteros
    return images, masks

# Cargar los datos
X_train, y_train = load_images_and_masks(train_image_dir,
train_mask_dir, IMAGE_SIZE)
X_val, y_val = load_images_and_masks(val_image_dir, val_mask_dir,
IMAGE_SIZE)

# Convertir máscaras a categóricas (One-Hot Encoding)
y_train = tf.keras.utils.to_categorical(y_train,
num_classes=NUM_CLASSES)
y_val = tf.keras.utils.to_categorical(y_val,
num_classes=NUM_CLASSES)

print(f"Imágenes de entrenamiento: {X_train.shape}, Máscaras:
{y_train.shape}")
print(f"Imágenes de validación: {X_val.shape}, Máscaras:
{y_val.shape}")

def weighted_loss(y_true, y_pred):
    class_weights = tf.constant([0.2, 22.98, 26.25, 25.42],
dtype=tf.float32) # Pesos de ejemplo
    y_true = tf.cast(tf.argmax(y_true, axis=-1), tf.int32)
    sample_weights = tf.gather(class_weights, y_true)
    cce = CategoricalCrossentropy(from_logits=False)
    loss = cce(y_true=tf.one_hot(y_true, depth=4), y_pred=y_pred)
    return loss * sample_weights

# Definir el modelo U-Net
def unet_model(input_size=(256, 256, 3), num_classes=4):
    inputs = Input(input_size)
    s = tf.keras.layers.Lambda(lambda x: x / 255.0)(inputs) #
Normalización integrada

    # Encoder
    c1 = Conv2D(64, (3, 3), activation='relu',
kernel_initializer='he_normal',
padding='same', kernel_regularizer=l2(0.001))(s)
    c1 = Dropout(0.4)(c1)
    c1 = Conv2D(64, (3, 3), activation='relu',
kernel_initializer='he_normal',
padding='same', kernel_regularizer=l2(0.001))(c1)
    p1 = MaxPooling2D((2, 2))(c1)

```

```

        c2 = Conv2D(128, (3, 3), activation='relu',
kernel_initializer='he_normal',
                    padding='same', kernel_regularizer=l2(0.001))(p1)
        c2 = Dropout(0.4)(c2)
        c2 = Conv2D(128, (3, 3), activation='relu',
kernel_initializer='he_normal',
                    padding='same', kernel_regularizer=l2(0.001))(c2)
        p2 = MaxPooling2D((2, 2))(c2)

        # Bottleneck
        c3 = Conv2D(256, (3, 3), activation='relu',
kernel_initializer='he_normal',
                    padding='same', kernel_regularizer=l2(0.001))(p2)
        c3 = Dropout(0.4)(c3)
        c3 = Conv2D(256, (3, 3), activation='relu',
kernel_initializer='he_normal',
                    padding='same', kernel_regularizer=l2(0.001))(c3)

        # Decoder
        u1 = UpSampling2D((2, 2))(c3)
        u1 = Concatenate()([u1, c2])
        c4 = Conv2D(128, (3, 3), activation='relu',
kernel_initializer='he_normal',
                    padding='same')(u1)
        c4 = Dropout(0.4)(c4)
        c4 = Conv2D(128, (3, 3), activation='relu',
kernel_initializer='he_normal',
                    padding='same')(c4)

        u2 = UpSampling2D((2, 2))(c4)
        u2 = Concatenate()([u2, c1])
        c5 = Conv2D(64, (3, 3), activation='relu',
kernel_initializer='he_normal',
                    padding='same')(u2)
        c5 = Dropout(0.4)(c5)
        c5 = Conv2D(64, (3, 3), activation='relu',
kernel_initializer='he_normal',
                    padding='same')(c5)

        outputs = Conv2D(num_classes, (1, 1),
activation='softmax')(c5)

        model = Model(inputs, outputs)
        return model

    # Crear el modelo
    model = unet_model(input_size=(IMAGE_SIZE[0], IMAGE_SIZE[1], 3),
num_classes=NUM_CLASSES)
    model.summary()

    # Compilar el modelo
    model.compile(optimizer=Adam(learning_rate=1e-4),
                    loss=weighted_loss,
                    metrics=['accuracy',
tf.keras.metrics.MeanIoU(num_classes=NUM_CLASSES), Precision(),
Recall())])

```

```

# Configurar el callback para guardar el modelo con los mejores
pesos
checkpoint = ModelCheckpoint(
    filepath=r"C:\Users\ismae\Desktop\TFG\model_unet_best3_1.h5",
    # Guarda el modelo completo con los mejores pesos
    monitor='val_loss', # Métrica para determinar los mejores
pesos
mejores
    save_best_only=True, # Solo guardar cuando los pesos son
    save_weights_only=False, # Guardar el modelo completo
    mode='min', # Buscar la pérdida mínima
    verbose=1
)
early_stopping = EarlyStopping(monitor='val_loss', patience=10,
restore_best_weights=True)

# Entrenar el modelo
history = model.fit(
    X_train, y_train,
    validation_data=(X_val, y_val),
    batch_size=BATCH_SIZE,
    epochs=EPOCHS,
    verbose=1,
    callbacks=[checkpoint, early_stopping] # Callback para
guardar los mejores pesos
)

# Guardar el modelo completo al final del entrenamiento
model.save(r"C:\Users\ismae\Desktop\TFG\model_unet_final3_1.h5")

# Graficar la historia de entrenamiento
plt.figure(figsize=(12, 6))
plt.plot(history.history['accuracy'], label='Precisión de
entrenamiento')
plt.plot(history.history['val_accuracy'], label='Precisión de
validación')
plt.plot(history.history['loss'], label='Pérdida de
entrenamiento')
plt.plot(history.history['val_loss'], label='Pérdida de
validación')
plt.legend()
plt.title("Historia de entrenamiento")
plt.show()

```